



# Machine Learning

## - Final project -


EGEI 2023

Iulia Florea  
Guido Pistorius  
Jairo Ramirez Torres



# Table of contents

- 01 → Data selection and cleaning
- 02 → Data Analysis and Variable selection
- 03 → ML Models: Logistic, KNN,  
Decision Tree, Random Forest
- 04 → K-means clustering
- 05 → Conclusion



01 →

# Data selection and cleaning

- Dataset: Heart Failure Prediction
- Variables: 5 numerical and 7 categorical
  - Target variable: HeartDisease (0 or 1)

	Age	Sex	ChestPainType	RestingBP	Cholesterol	FastingBS	RestingECG		MaxHR	ExerciseAngina	Oldpeak	ST_Slope	HeartDisease
0	40	M	ATA	140	289	0	Normal	0	172	N	0.0	Up	0
1	49	F	NAP	160	180	0	Normal	1	156	N	1.0	Flat	1
2	37	M	ATA	130	283	0	ST	2	98	N	0.0	Up	0
3	48	F	ASY	138	214	0	Normal	3	108	Y	1.5	Flat	1
4	54	M	NAP	150	195	0	Normal	4	122	N	0.0	Up	0
...	...	...	...	...	...	...	...	...	...	...	...	...	...
913	45	M	TA	110	264	0	Normal	913	132	N	1.2	Flat	1
914	68	M	ASY	144	193	1	Normal	914	141	N	3.4	Flat	1
915	57	M	ASY	130	131	0	Normal	915	115	Y	1.2	Flat	1
916	57	F	ATA	130	236	0	LVH	916	174	N	0.0	Flat	1
917	38	M	NAP	138	175	0	Normal	917	173	N	0.0	Up	0

Missing values

Age	0.0
Sex	0.0
ChestPainType	0.0
RestingBP	0.0
Cholesterol	0.0
FastingBS	0.0
RestingECG	0.0
MaxHR	0.0
ExerciseAngina	0.0
Oldpeak	0.0
ST_Slope	0.0
HeartDisease	0.0

- Data cleaning
  - Encoding categorical variables
  - Checking for missing values
  - Converting the values to *float* type

```
#read the data
data = pd.read_csv('heart.csv')

# Use the map function to replace categorical values
data['Sex'] = data['Sex'].map({'M': 1, 'F': 0})
data['ChestPainType'] = data['ChestPainType'].map({'TA': 1, 'ATA': 2, 'NAP':3, 'ASY':4})
data['RestingECG'] = data['RestingECG'].map({'Normal':1, 'ST':2, 'LVH':3})
data['ExerciseAngina'] = data['ExerciseAngina'].map({'Y':1, 'N':0})
data['ST_Slope'] = data['ST_Slope'].map({'Down':1, 'Flat':2, 'Up':3})

# check number of missing values
missing_values = data.isnull()
print(missing_values.sum()/len(data))# There are no missing values in the data

#convert to float type
data = data.astype(float)
print(data)
```

# Data after encoding

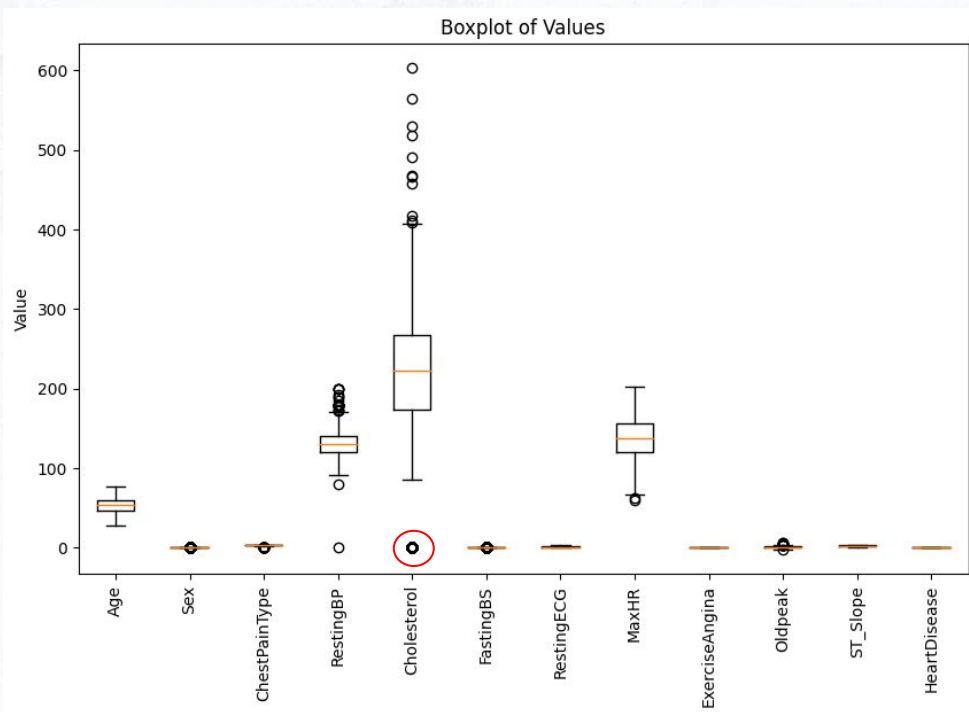
```
dtype: float64
Age Sex ChestPainType RestingBP Cholesterol FastingBS RestingECG 0 MaxHR ExerciseAngina Oldpeak ST_Slope HeartDisease
0 40.0 1.0 2.0 140.0 289.0 0.0 1.0 1 156.0 0.0 1.0 2.0 1.0
1 49.0 0.0 3.0 160.0 180.0 0.0 1.0 2 98.0 0.0 0.0 3.0 0.0
2 37.0 1.0 2.0 130.0 283.0 0.0 2.0 3 108.0 1.0 1.5 2.0 1.0
3 48.0 0.0 4.0 138.0 214.0 0.0 1.0 4 122.0 0.0 0.0 3.0 0.0
4 54.0 1.0 3.0 150.0 195.0 0.0 1.0 .. ... ... ...
.. ... ... ... 913 132.0 0.0 1.2 2.0 1.0
913 45.0 1.0 1.0 110.0 264.0 0.0 1.0 914 141.0 0.0 3.4 2.0 1.0
914 68.0 1.0 4.0 144.0 193.0 1.0 1.0 915 115.0 1.0 1.2 2.0 1.0
915 57.0 1.0 4.0 130.0 131.0 0.0 1.0 916 174.0 0.0 0.0 2.0 1.0
916 57.0 0.0 2.0 130.0 236.0 0.0 3.0 917 173.0 0.0 0.0 3.0 0.0
917 38.0 1.0 3.0 138.0 175.0 0.0 1.0
[918 rows x 12 columns]
```




02 →

# Data analysis and Variable selection


# Data analysis




- Boxplot to observe outliers
- The boxplot shows the value 0 for cholesterol, which is not possible → Missing values



```
plt.figure(figsize=(10, 6))
plt.boxplot(data)
plt.title('Boxplot of Values')
plt.ylabel('Value')
plt.xticks(list(range(1,13)),list(data.columns),rotation = 90)
plt.show()
```



Cholesterol histogram



Cholesterol histogram without 0 values


- The histogram shows how the values for cholesterol are distributed  
→ many 0 values
- 0 values = missing values → biased results
- We eliminate the observations that have the value 0 for cholesterol

```
plt.hist(data['Cholesterol'])  
data = data[data['Cholesterol'] != 0]  
plt.hist(data['Cholesterol'], color='orange')
```



# Removing outliers

Boxplot of Values After Filtering



- We filtered the data to remove the outliers for the variables:
  - Age
  - Cholesterol
  - Resting Blood Pressure (Resting BP)
  - Maximum Heart Rate Achieved (MaxHR)
- After removing the outliers we have 704 observations

```
plt.figure(figsize=(10, 6))
plt.boxplot(data)
plt.title('Boxplot of Values')
plt.ylabel('Value')
plt.xticks(list(range(1,13)),list(data.columns),rotation = 90)
plt.show()
```

## More data cleaning

- Eliminate the OldPeak variable, since some models do not accept negative values
- Select the target variable: HeartDisease (0 or 1)

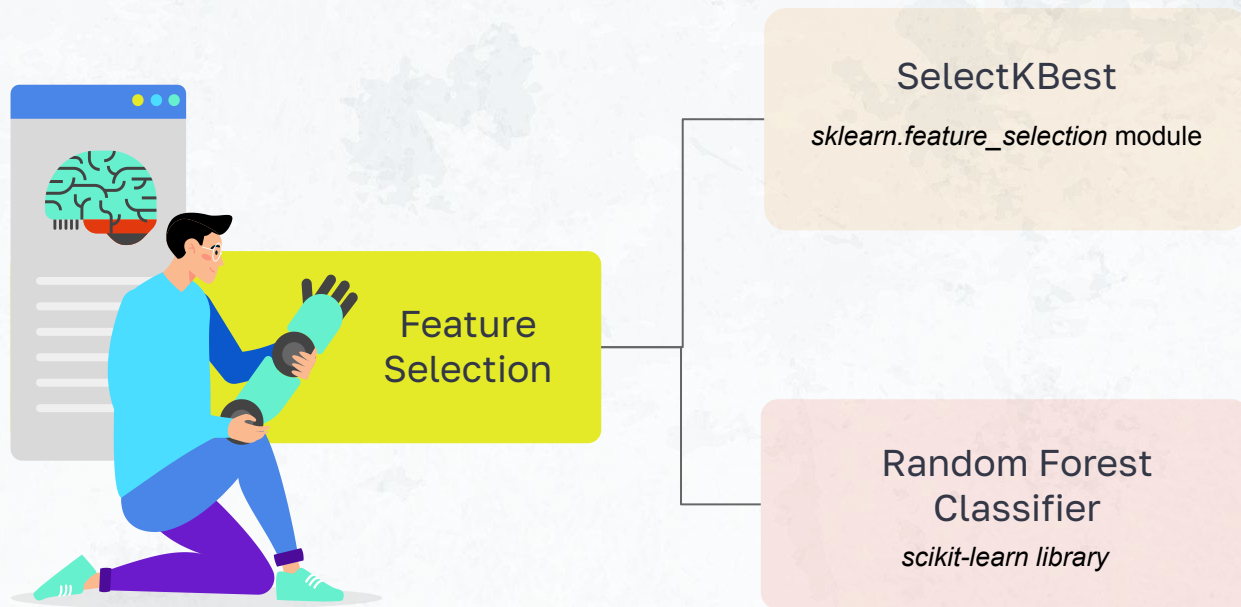
```
data_filtered = data_filtered.drop(['Oldpeak'], axis = 1)

# Extracting labels and saving it in y
y = data_filtered['HeartDisease']
# Dropping labels from dataset
X = data_filtered.drop(['HeartDisease'],axis=1)
```




# Feature selection

- Choosing the most relevant features for our models is important, as they will contribute to the predictive power
- We are using 2 methods to identify the features and we combine the top 2 from each of them, to be used in the models



Feature Importances in the Heart Dataset



# Select KBest Model

Feature ranking:


1. feature MaxHR (463.873221) SELECTED
2. feature ExerciseAngina (134.040787) SELECTED
3. feature Age (120.092089) SELECTED
4. feature Cholesterol (61.611139) SELECTED
5. feature ChestPainType (44.891665)
6. feature ST\_Slope (37.145919)
7. feature RestingBP (37.088694)
8. feature Sex (15.978849)
9. feature FastingBS (15.231953)
10. feature RestingECG (4.440008)

```
# Selecting the top 2 features
selector = SelectKBest(score_func=chi2, k = 3)
X_selected = selector.fit_transform(X, y)

# Get the indices that would sort the scores in descending order
indices = np.argsort(selector.scores_)[::-1]
# Displaying the feature importances
print("Feature ranking:")
for rank, index in enumerate(indices):
    if rank < 4: print("%d. feature %s (%f) SELECTED" % (rank + 1, X.columns[index], selector.scores_[index]))
    else: print("%d. feature %s (%f)" % (rank + 1, X.columns[index], selector.scores_[index]))
```



Feature Importances in the Heart Dataset



# Random Forest Classifier

Feature ranking:

1. feature ST\_Slope (0.260708)
2. feature ChestPainType (0.129920)
3. feature MaxHR (0.125942)
4. feature Age (0.112269)
5. feature ExerciseAngina (0.111505)
6. feature Cholesterol (0.088764)
7. feature RestingBP (0.077132)
8. feature Sex (0.049424)
9. feature RestingECG (0.032864)
10. feature FastingBS (0.011471)

```
rf = RandomForestClassifier(n_estimators=100)
rf.fit(X, y)

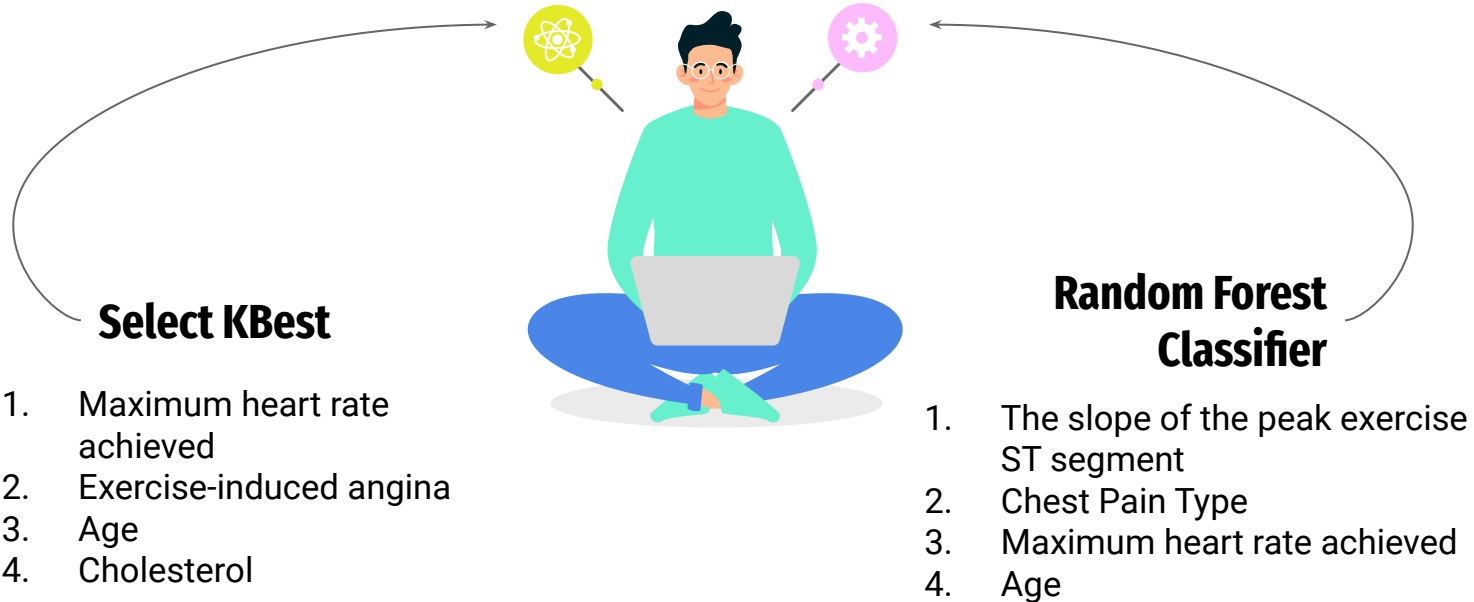
# Getting feature importances
importances = rf.feature_importances_
# Sorting the feature importances in descending order
indices = np.argsort(importances)[::-1]
# Displaying the feature importances
print("Feature ranking:")
for f in range(X.shape[1]):
    print("%d. feature %s (%f)" % (f + 1, X.columns[indices[f]], importances[indices[f]]))
```



# Final Feature Selection



1. Maximum heart rate achieved
2. Exercise-induced angina
3. The slope of the peak exercise ST segment
4. Chest Pain Type



03 →

**ML Models: Logistic, KNN,  
Decision Tree, Random Forest**

# Types of models

## (a) Logistic Regression

Logistic Regression is a statistical model used for binary classification that estimates the probability of an instance belonging to a **particular class based on a linear combination** of input features.

## (b) K-Nearest Neighbors (KNN)


K-Nearest Neighbors is a non-parametric and instance-based machine learning algorithm that **classifies a data point based on the majority class of its k nearest neighbors** in the feature space.

## (c) Decision Tree

A Decision Tree is a **tree-like model** where **internal nodes** represent **decisions** based on features, branches represent possible outcomes, and leaf nodes represent the final classification or regression result.

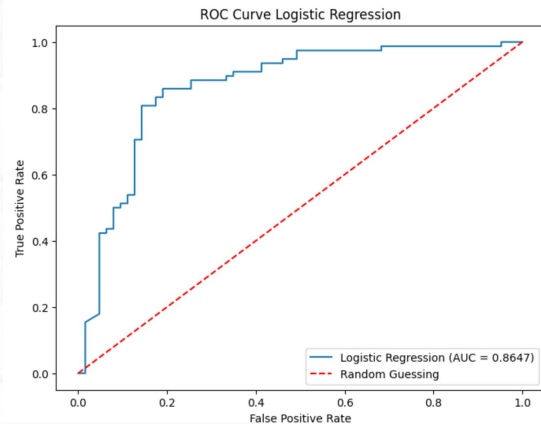
## (d) Random Forest

Random Forest is an ensemble learning method that builds **multiple decision trees** during training and outputs the mode of the classes (classification) or the mean prediction (regression) of the individual trees as the final result.

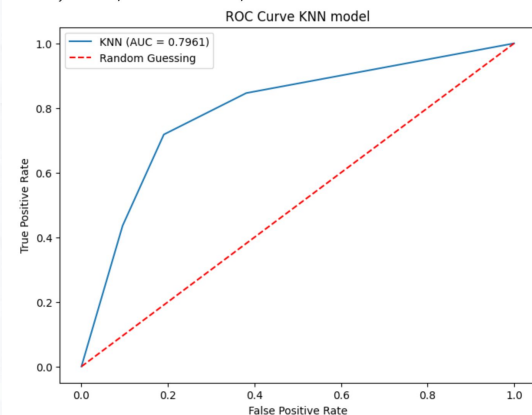


# Logistic regression & KNN

Logistic Regression Metrics:  
Accuracy: 0.8227, Precision: 0.8732, Recall: 0.7949



KNN Metrics:  
Accuracy: 0.7589, Precision: 0.8235, Recall: 0.7179



```
# Splitting data into training and testing sets
```

```
X_train, X_test, y_train, y_test = train_test_split(X_sel, y, test_size=0.2, random_state=45)
```

```
# Initialize and train logistic regression model
```

```
logreg = LogisticRegression()  
logreg.fit(X_train, y_train)
```

```
# Predictions and Evaluation for Logistic Regression
```

```
y_pred_logreg = logreg.predict(X_test)
```

```
# Calculating metrics for Logistic Regression
```

```
accuracy_logreg = accuracy_score(y_test, y_pred_logreg)  
precision_logreg = precision_score(y_test, y_pred_logreg)  
recall_logreg = recall_score(y_test, y_pred_logreg)
```

```
# ROC Curve for Logistic Regression
```

```
y_proba_logreg = logreg.predict_proba(X_test)[: , 1]  
fpr_logreg, tpr_logreg, thresholds_logreg = roc_curve(y_test, y_proba_logreg)  
roc_auc_logreg = roc_auc_score(y_test, y_proba_logreg)
```

```
# Splitting data into training and testing sets
```

```
X_train, X_test, y_train, y_test = train_test_split(X_sel, y, test_size=0.2, random_state=45)
```

```
# Using KNN for classification with k=4
```

```
knn = KNeighborsClassifier(n_neighbors=3)  
knn.fit(X_train, y_train)
```

```
# Make the predictions
```

```
y_pred_knn = knn.predict(X_test)
```

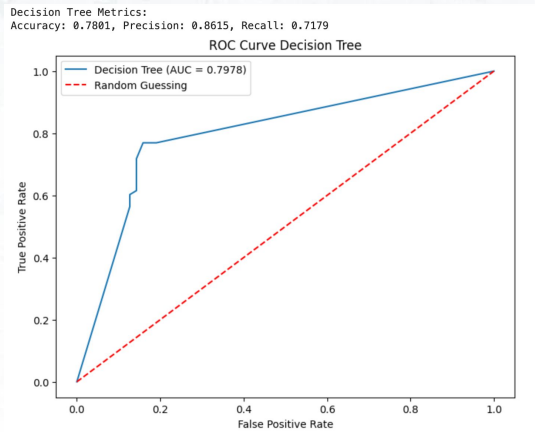
```
# Calculating metrics for KNN
```

```
accuracy_knn = accuracy_score(y_test, y_pred_knn)  
precision_knn = precision_score(y_test, y_pred_knn)  
recall_knn = recall_score(y_test, y_pred_knn)
```

```
# ROC Curve for KNN
```

```
y_proba_knn = knn.predict_proba(X_test)[: , 1]  
fpr_knn, tpr_knn, thresholds_knn = roc_curve(y_test, y_proba_knn)  
roc_auc_knn = roc_auc_score(y_test, y_proba_knn)
```

# Decision Tree & Random Forest

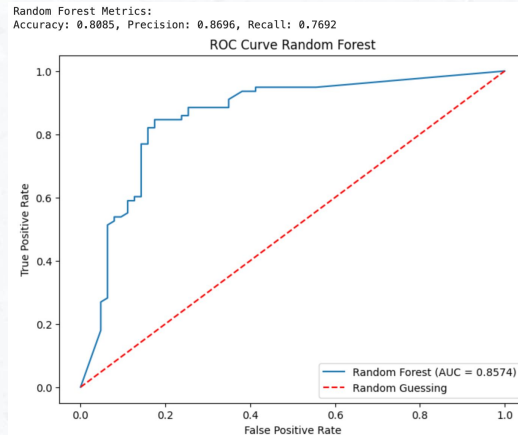


```
# Initialize and train the decision tree classifier
clf = DecisionTreeClassifier(random_state=42)
clf.fit(X_train, y_train)

# Predictions and Evaluation for Decision Tree
y_pred_dt = clf.predict(X_test)

# Calculating metrics for Decision Tree
accuracy_dt = accuracy_score(y_test, y_pred_dt)
precision_dt = precision_score(y_test, y_pred_dt)
recall_dt = recall_score(y_test, y_pred_dt)

# ROC Curve for Decision Tree
y_proba_dt = clf.predict_proba(X_test)[:, 1]
fpr_dt, tpr_dt, thresholds_dt = roc_curve(y_test, y_proba_dt)
roc_auc_dt = roc_auc_score(y_test, y_proba_dt)
```



```
# Initialize and train the Random Forest classifier
rf_clf = RandomForestClassifier(random_state=42)
rf_clf.fit(X_train, y_train)

# Predictions and Evaluation for Random Forest
y_pred_rf = rf_clf.predict(X_test)

# Calculating metrics for Random Forest
accuracy_rf = accuracy_score(y_test, y_pred_rf)
precision_rf = precision_score(y_test, y_pred_rf)
recall_rf = recall_score(y_test, y_pred_rf)

# ROC Curve for Random Forest
y_proba_rf = rf_clf.predict_proba(X_test)[:, 1]
fpr_rf, tpr_rf, thresholds_rf = roc_curve(y_test, y_proba_rf)
roc_auc_rf = roc_auc_score(y_test, y_proba_rf)
```



# Final Comparison Table

<b>Model</b>	<i>Accuracy</i>	<i>Precision</i>	<i>Recall</i>	<i>ROC-AUC</i>
<i>Logistic regression</i>	<b>0.8227</b>	<b>0.8732</b>	<b>0.7949</b>	<b>0.8647</b>
<i>KNN</i>	0.7801	0.8235	0.7179	0.7961
<i>Decision Tree</i>	0.7801	0.8615	0.7179	0.7978
<i>Random Forest</i>	0.8085	0.8696	0.7692	0.8574

- The **logistic regression** outperforms the other models taking into account all the test statistics

04 →

# K-means clustering

# Algorithm

```
scaler = StandardScaler()
X[['Age', 'Cholesterol', 'RestingBP', 'MaxHR']] =
    scaler.fit_transform(X[['Age', 'Cholesterol', 'RestingBP', 'MaxHR']])

kmeans = KMeans(
    init="random",
    n_clusters=2,
    n_init=10,
    max_iter=300,
    random_state=41
)



kmeans.fit(X)
print(kmeans.inertia_)
print(kmeans.cluster_centers_)
print(kmeans.n_iter_)
print(kmeans.labels_)
X['clusters'] = kmeans.labels_

kmeans_kwargs = {
    "init": "random",
    "n_init": 10,
    "max_iter": 300,
    "random_state": 42
}
```

- Step 1: Choose number of K
- Step 2: Randomly initialize centroids
- Steps 3-4: Iterative process called expectation-maximization
  - Expectation assigns each point to closest centroid
  - Maximization computes the new centroid mean of each cluster
  - Repeat until centroid positions do not change

# Choosing K

- To choose the optimal numbers of K, we used Silhouette and Elbow Method.
- In both cases, K=2 is optimal



# Cluster analysis

Cluster 0 is composed of 295 men and 58 women that are old on average, with higher than average Resting BP, Cholesterol, Fasting BS, Resting ECG, and Exercise Angina leading to large proportion of them having heart disease.

Cluster 1, on the other hand, are a younger group of 243 men and 108 women with lower than average Resting BP, Cholesterol, Fasting BS, Resting ECG, and Exercise Angina leading to a smaller proportion of them having heart disease.

Mean	Age	Sex	Chest Pain Type	Resting BP	Cholesterol	Fasting BS	Resting ECG	Max HR	Exercise Angina	ST Slope	Heart Disease
X (704)	0	0.764	3.149	0	0	0.161	1.634	0	0.381	2.420	<b>0.472</b>
Cluster 0 (353)	0.507	0.836	3.737	0.238	0.080	0.227	1.782	-0.665	0.657	2.108	<b>0.751</b>
Cluster 1 (351)	-0.509	0.692	2.558	-0.239	-0.080	0.094	1.484	0.669	0.103	2.735	<b>0.191</b>



# Thank you!

---

